

benchmark_k6_read

July 13, 2025

```
[25]: import json
import glob
import re
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Suffix for selecting benchmark results
suffix = "/read_avgnet/"

# Path to your benchmark results directory
results_dir = './results_store/_final/'
ec_pattern = results_dir + f'erasure{suffix}/_read_*.json'
replication_pattern = results_dir + f'replication{suffix}/_read_*.json'

# Helper to extract data from files

# Helper to extract data from files
def extract_data(files, system_type):
    data = []
    for file in files:
        # Updated regex: matches kbit, mbit, gbit
        match = re.search(r'_read_(.+?)_(\d+)vu(?:_(?:[0-9]+(?:?:kbit|mbit|gbit)))?↪\.json', file)
        if not match:
            continue
        payload_size = match.group(1)
        if payload_size.endswith('b'):
            payload_size = payload_size[:-1]
        try:
            payload_size = int(payload_size)
        except ValueError:
            pass
        virtual_user = match.group(2)
        bandwidth = match.group(3) if match.group(3) else 'unlimited'
        with open(file) as f:
            j = json.load(f)
```

```

        # try summary.success_performance, else fallback to details.
        ↪http_req_duration
        sp = j.get('summary', {}).get('success_performance')
        if sp:
            med = sp.get('med', 0)
            p90 = sp.get('p(90)', 0)
            avg = sp.get('avg', 0)
        else:
            dur = j.get('details', {}).get('http_req_duration', {}).
            ↪get('values', {})
            med = dur.get('med', 0)
            p90 = dur.get('p(90)', 0)
            avg = dur.get('avg', 0)
        # extract request rate
        rate = j.get('summary', {}).get('reqs', {}).get('rate', 0)
        data.append({
            'system': system_type,
            'payload_size': payload_size,
            'virtual_user': int(virtual_user),
            'bandwidth': bandwidth,
            'med': med,
            'p90': p90,
            'avg': avg,
            'rate': rate
        })
    return data

def bandwidth_to_num(bw):
    if bw == 'unlimited':
        return float('inf')
    m = re.match(r'(\d+)(kbit|mbit|gbit)', bw)
    if not m:
        return float('inf')
    val, unit = int(m.group(1)), m.group(2)
    if unit == 'kbit':
        return val
    elif unit == 'mbit':
        return val * 1000
    elif unit == 'gbit':
        return val * 1000 * 1000
    return float('inf')

# Collect EC and Replication data
files_ec = glob.glob(ec_pattern)
files_replication = glob.glob(replication_pattern)
data = extract_data(files_ec, 'EC') + extract_data(files_replication,
↪'Replication')

```

```

df = pd.DataFrame(data)
if not df.empty:
    # Add bandwidth_num for correct sorting
    df['bandwidth_num'] = df['bandwidth'].apply(bandwidth_to_num)
    # Sort by payload_size, virtual_user, bandwidth_num, then system for
    ↪grouped comparison
    df = df.sort_values(['payload_size', 'virtual_user', 'bandwidth_num',
    ↪'system'])
    df['combo'] = df.apply(lambda row:
    ↪f"{row['payload_size']}B_{row['virtual_user']}vu_{row['bandwidth']}", axis=1)
    df.reset_index(drop=True, inplace=True)
df

```

```

[25]:
      system  payload_size  virtual_user  bandwidth      med      p90  \
0         EC      200000           1    10mbit   9.124901  10.055491
1  Replication      200000           1    10mbit   2.551200   3.139400
2         EC      200000           1    25mbit   8.878701  10.085901
3  Replication      200000           1    25mbit   2.214601   2.711220
4         EC      200000           1    40mbit   8.708501   9.745851
5  Replication      200000           1    40mbit   2.206701   2.701601
6         EC      200000           1    55mbit   8.527501   9.542181
7  Replication      200000           1    55mbit   2.203500   2.690460
8         EC      200000           1    70mbit   8.179301   9.215161
9  Replication      200000           1    70mbit   2.144900   2.622800
10        EC      400000           1    10mbit  12.772001  16.404281
11  Replication      400000           1    10mbit   3.159750   4.113140
12        EC      400000           1    25mbit  12.520901  15.096402
13  Replication      400000           1    25mbit   2.842400   3.647481
14        EC      400000           1    40mbit  12.175551  13.843451
15  Replication      400000           1    40mbit   2.858900   3.801000
16        EC      400000           1    55mbit  11.962101  13.686821
17  Replication      400000           1    55mbit   2.866700   3.818521
18        EC      400000           1    70mbit  11.404601  12.973821
19  Replication      400000           1    70mbit   2.743200   3.581411
20        EC      600000           1    10mbit  15.724501  17.259242
21  Replication      600000           1    10mbit   3.845950   4.728531
22        EC      600000           1    25mbit  15.408301  17.127521
23  Replication      600000           1    25mbit   3.260801   3.969001
24        EC      600000           1    40mbit  15.225651  16.935872
25  Replication      600000           1    40mbit   3.250301   3.964220
26        EC      600000           1    55mbit  14.982252  16.750471
27  Replication      600000           1    55mbit   3.226001   3.908161
28        EC      600000           1    70mbit  14.340301  16.054062
29  Replication      600000           1    70mbit   3.196500   3.886910
30        EC      800000           1    10mbit  19.107251  21.944972
31  Replication      800000           1    10mbit   4.211951   5.957451
32        EC      800000           1    25mbit  18.645702  21.561252

```

33	Replication	800000	1	25mbit	3.962301	6.422200
34	EC	800000	1	40mbit	18.405901	20.948841
35	Replication	800000	1	40mbit	3.885200	5.902941
36	EC	800000	1	55mbit	18.267102	22.380782
37	Replication	800000	1	55mbit	3.894900	6.102000
38	EC	800000	1	70mbit	17.485401	21.007102
39	Replication	800000	1	70mbit	3.812800	5.708961
40	EC	1000000	1	10mbit	22.146502	25.057702
41	Replication	1000000	1	10mbit	4.604700	6.251681
42	EC	1000000	1	25mbit	21.368901	24.712341
43	Replication	1000000	1	25mbit	4.214101	5.506330
44	EC	1000000	1	40mbit	21.237251	24.414312
45	Replication	1000000	1	40mbit	4.214251	5.480670
46	EC	1000000	1	55mbit	20.939201	23.878402
47	Replication	1000000	1	55mbit	4.298851	6.445801
48	EC	1000000	1	70mbit	20.203501	23.103322
49	Replication	1000000	1	70mbit	4.151101	5.356040

	avg	rate	bandwidth_num	combo
0	10.860901	71.560057	10000	200000B_1vu_10mbit
1	2.985204	172.481319	10000	200000B_1vu_10mbit
2	27.179235	32.630497	25000	200000B_1vu_25mbit
3	2.354469	195.552406	25000	200000B_1vu_25mbit
4	11.303761	69.505690	40000	200000B_1vu_40mbit
5	2.305885	197.046616	40000	200000B_1vu_40mbit
6	10.762659	72.170053	55000	200000B_1vu_55mbit
7	2.293517	198.336889	55000	200000B_1vu_55mbit
8	14.392082	56.709423	70000	200000B_1vu_70mbit
9	2.230576	200.676736	70000	200000B_1vu_70mbit
10	67.232828	13.505693	10000	400000B_1vu_10mbit
11	4.684858	101.690860	10000	400000B_1vu_10mbit
12	62.172990	14.587000	25000	400000B_1vu_25mbit
13	3.395078	119.415664	25000	400000B_1vu_25mbit
14	15.406891	47.675817	40000	400000B_1vu_40mbit
15	3.287103	119.124967	40000	400000B_1vu_40mbit
16	14.826273	48.919919	55000	400000B_1vu_55mbit
17	3.187288	120.100178	55000	400000B_1vu_55mbit
18	14.857108	49.250975	70000	400000B_1vu_70mbit
19	2.989706	125.451981	70000	400000B_1vu_70mbit
20	20.531078	36.227728	10000	600000B_1vu_10mbit
21	9.238988	60.157541	10000	600000B_1vu_10mbit
22	18.904356	38.453432	25000	600000B_1vu_25mbit
23	4.245174	91.746041	25000	600000B_1vu_25mbit
24	18.595093	38.904588	40000	600000B_1vu_40mbit
25	3.846926	95.439353	40000	600000B_1vu_40mbit
26	18.285675	39.244829	55000	600000B_1vu_55mbit
27	3.627947	97.371060	55000	600000B_1vu_55mbit

28	18.862731	38.363590	70000	600000B_1vu_70mbit
29	3.511334	97.896048	70000	600000B_1vu_70mbit
30	27.075902	26.900681	10000	800000B_1vu_10mbit
31	17.297442	37.215449	10000	800000B_1vu_10mbit
32	24.437212	29.233778	25000	800000B_1vu_25mbit
33	6.324088	62.585153	25000	800000B_1vu_25mbit
34	23.716280	29.908257	40000	800000B_1vu_40mbit
35	5.348494	68.060497	40000	800000B_1vu_40mbit
36	23.316528	29.731001	55000	800000B_1vu_55mbit
37	5.087205	68.867878	55000	800000B_1vu_55mbit
38	22.989338	30.217363	70000	800000B_1vu_70mbit
39	4.749964	70.600655	70000	800000B_1vu_70mbit
40	33.255462	22.189536	10000	1000000B_1vu_10mbit
41	52.175193	15.402991	10000	1000000B_1vu_10mbit
42	29.058840	24.496383	25000	1000000B_1vu_25mbit
43	6.916651	56.240790	25000	1000000B_1vu_25mbit
44	27.944875	25.200413	40000	1000000B_1vu_40mbit
45	5.903930	60.001215	40000	1000000B_1vu_40mbit
46	26.359037	26.376127	55000	1000000B_1vu_55mbit
47	5.739029	59.012397	55000	1000000B_1vu_55mbit
48	25.816443	26.766462	70000	1000000B_1vu_70mbit
49	5.126644	62.710817	70000	1000000B_1vu_70mbit

```
[26]: # Plotting: Compare EC vs Replication for each metric, grouped by
      ↪(payload_size, virtual_user)
sns.set_style("whitegrid")
metrics = ['med', 'p90', 'avg', 'rate']
metric_titles = {
    'med': 'Median Latency (ms)',
    'p90': 'P90 Latency (ms)',
    'avg': 'Average Latency (ms)',
    'rate': 'Request Rate (req/s)'
}

fig, axes = plt.subplots(2, 2, figsize=(20, 12))
axes = axes.flatten()

for ax, metric in zip(axes, metrics):
    sns.barplot(
        data=df,
        x='combo',
        y=metric,
        hue='system',
        ci=None,
        dodge=True,
        ax=ax
    )
```

```

ax.set_title(f'EC vs Replication: {metric_titles[metric]}')
ax.set_xlabel('Payload Size & Virtual Users')
ax.set_ylabel(metric_titles[metric])
ax.tick_params(axis='x', rotation=90)

# single legend for all subplots
handles, labels = axes[0].get_legend_handles_labels()
fig.legend(handles, labels, title='System', loc='upper right')

fig.tight_layout()
plt.show()

```

/tmp/ipykernel_288815/1184367366.py:15: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```

sns.barplot(
/tmp/ipykernel_288815/1184367366.py:15: FutureWarning:

```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```

sns.barplot(
/tmp/ipykernel_288815/1184367366.py:15: FutureWarning:

```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```

sns.barplot(
/tmp/ipykernel_288815/1184367366.py:15: FutureWarning:

```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```

sns.barplot(

```



[27]: *# Create grouped visualizations for bandwidth, payload size, and virtual users*

```
# First, let's check what distinct values we have for each parameter
print(f"Unique bandwidth values: {df['bandwidth'].unique()}")
print(f"Unique payload sizes: {df['payload_size'].unique()}")
print(f"Unique virtual users: {df['virtual_user'].unique()}")
```

```
Unique bandwidth values: ['10mbit' '25mbit' '40mbit' '55mbit' '70mbit']
Unique payload sizes: [ 200000  400000  600000  800000 1000000]
Unique virtual users: [1]
```

[28]: *# 1. Grouping by bandwidth*

```
def plot_by_bandwidth():
    # Sort bandwidths numerically using bandwidth_num
    bw_order = (
        df[['bandwidth', 'bandwidth_num']]
        .drop_duplicates()
        .sort_values('bandwidth_num')
        .bandwidth.tolist()
    )
    for bw in bw_order:
        df_bw = df[df['bandwidth'] == bw]
        if df_bw.empty:
            continue

    # set up 2 x 2 subplots for the four metrics
```

```

fig, axes = plt.subplots(2, 2, figsize=(20, 12))
axes = axes.flatten()
for ax, metric in zip(axes, metrics):
    # label for x-axis
    df_bw['payload_vu'] = df_bw.apply(
        lambda row: f"{row['payload_size']}B_{row['virtual_user']}vu",
        axis=1
    )
    sns.barplot(
        data=df_bw,
        x='payload_vu',
        y=metric,
        hue='system',
        ci=None,
        dodge=True,
        ax=ax
    )
    ax.set_title(f'EC vs Replication: {metric_titles[metric]} at {bw} Bandwidth')
    ax.set_xlabel('Payload Size & Virtual Users')
    ax.set_ylabel(metric_titles[metric])
    ax.tick_params(axis='x', rotation=90)

    # single legend for all subplots
    handles, labels = axes[0].get_legend_handles_labels()
    fig.legend(handles, labels, title='System', loc='upper right')
    fig.suptitle(f'EC vs Replication Metrics at {bw} Bandwidth',
        fontsize=16)
    fig.tight_layout(rect=[0, 0, 1, 0.97])
    plt.show()

plot_by_bandwidth()

```

/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
```


Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(  

```



```
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
```

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.


```
sns.barplot(
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(
```



```
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  

```



```
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
```

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.


```
sns.barplot(
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(
```



```
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_288815/2927317137.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_288815/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  

```



```
[29]: # 2. Grouping by payload size (all metrics in one 2 x 2 display)
def plot_by_payload():
    payloads = df['payload_size'].unique()
    for payload in payloads:
        df_payload = df[df['payload_size'] == payload]
        if df_payload.empty:
            continue

        # Sort bandwidths numerically for this payload
        bw_order = (
            df_payload[['bandwidth', 'bandwidth_num']]
            .drop_duplicates()
            .sort_values('bandwidth_num')
            .bandwidth.tolist()
        )
        # Create combined x-axis labels and enforce ordering
        df_payload['bw_vu'] = df_payload.apply(
            lambda row: f"{row['bandwidth']}_{row['virtual_user']}_vu", axis=1
        )
        vu_order = sorted(df_payload['virtual_user'].unique())
        x_order = [f"{bw}_{vu}_vu" for bw in bw_order for vu in vu_order]
        df_payload['bw_vu'] = pd.Categorical(df_payload['bw_vu'],
        categories=x_order, ordered=True)

        # set up 2 x 2 subplots for the four metrics
```



```

fig, axes = plt.subplots(2, 2, figsize=(20, 12))
axes = axes.flatten()
for ax, metric in zip(axes, metrics):
    sns.barplot(
        data=df_payload,
        x='bw_vu',
        y=metric,
        hue='system',
        ci=None,
        dodge=True,
        ax=ax
    )
    ax.set_title(f"EC vs Replication: {metric_titles[metric]} for {payload}B")
    ax.set_xlabel("Bandwidth & Virtual Users")
    ax.set_ylabel(metric_titles[metric])
    ax.tick_params(axis='x', rotation=90)

    # single legend for all subplots
    handles, labels = axes[0].get_legend_handles_labels()
    fig.legend(handles, labels, title='System', loc='upper right')
    fig.suptitle(f"EC vs Replication Metrics for {payload}B Payload Size",
    ↪fontsize=16)
    fig.tight_layout(rect=[0, 0, 1, 0.96])
    plt.show()

plot_by_payload()

```

/tmp/ipykernel_288815/1240405451.py:17: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = df_payload.apply(
```

/tmp/ipykernel_288815/1240405451.py:22: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = pd.Categorical(df_payload['bw_vu'], categories=x_order,
ordered=True)
```

/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.


```
sns.barplot(
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
```



```
/tmp/ipykernel_288815/1240405451.py:17: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = df_payload.apply(
/tmp/ipykernel_288815/1240405451.py:22: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
```

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = pd.Categorical(df_payload['bw_vu'], categories=x_order,
ordered=True)
```

/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.


```
sns.barplot(
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(
```



```
/tmp/ipykernel_288815/1240405451.py:17: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = df_payload.apply(  
/tmp/ipykernel_288815/1240405451.py:22: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = pd.Categorical(df_payload['bw_vu'], categories=x_order,  
ordered=True)  
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

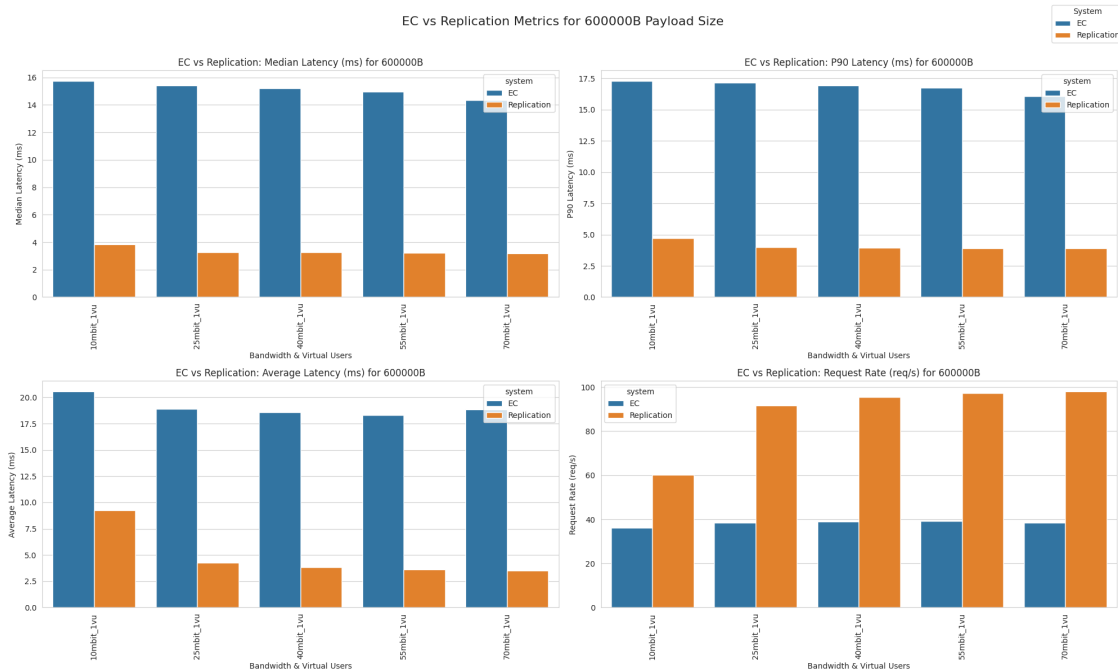
The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  

```



```
/tmp/ipykernel_288815/1240405451.py:17: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = df_payload.apply(
/tmp/ipykernel_288815/1240405451.py:22: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = pd.Categorical(df_payload['bw_vu'], categories=x_order,
ordered=True)
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

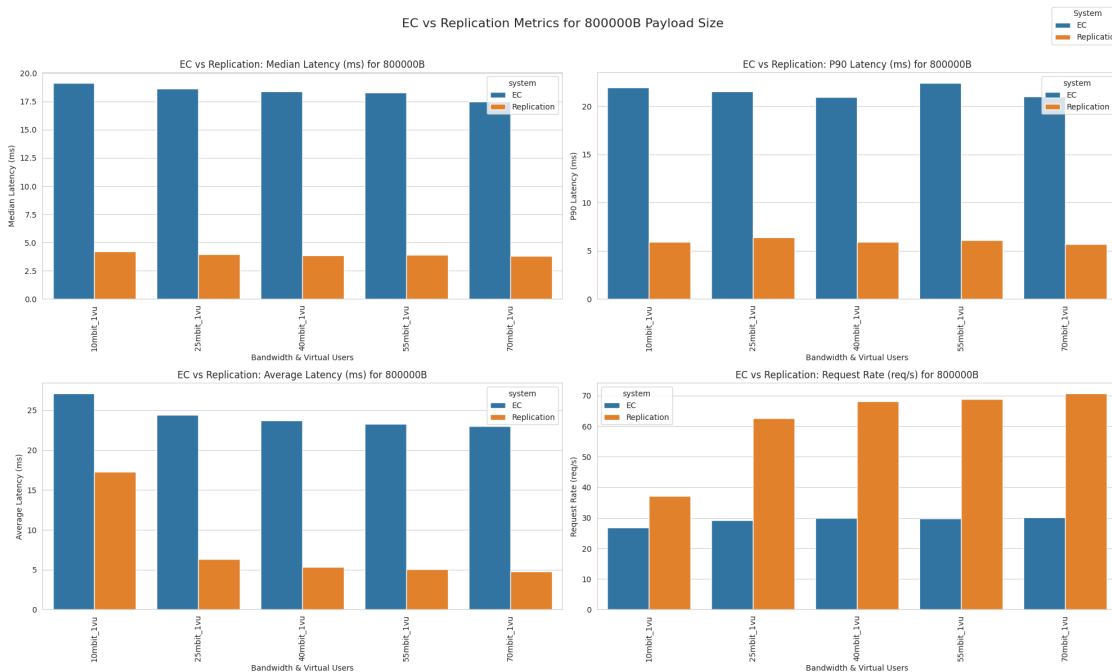
```
sns.barplot(
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
```



```
/tmp/ipykernel_288815/1240405451.py:17: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = df_payload.apply(
/tmp/ipykernel_288815/1240405451.py:22: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = pd.Categorical(df_payload['bw_vu'], categories=x_order,
ordered=True)
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.


```
sns.barplot(
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_288815/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
```



```
[30]: # 3. Grouping by virtual users (all metrics in one 2 x 2 display)
def plot_by_vu():
    vus = df['virtual_user'].unique()
    for vu in vus:
        df_vu = df[df['virtual_user'] == vu]
```

```

if df_vu.empty:
    continue

# set up 2 x 2 subplots for the four metrics
fig, axes = plt.subplots(2, 2, figsize=(20, 12))
axes = axes.flatten()

# precompute ordering for x-axis
bw_order = (
    df_vu[['bandwidth', 'bandwidth_num']]
    .drop_duplicates()
    .sort_values('bandwidth_num')
    .bandwidth.tolist()
)
ps_order = sorted(df_vu['payload_size'].unique())
x_order = [f"{bw}_{ps}B" for bw in bw_order for ps in ps_order]

# apply categorical ordering
df_vu['bw_payload'] = df_vu.apply(
    lambda row: f"{row['bandwidth']}_{row['payload_size']}B", axis=1
)
df_vu['bw_payload'] = pd.Categorical(df_vu['bw_payload'],
categories=x_order, ordered=True)

for ax, metric in zip(axes, metrics):
    sns.barplot(
        data=df_vu,
        x='bw_payload',
        y=metric,
        hue='system',
        ci=None,
        dodge=True,
        ax=ax
    )
    ax.set_title(f"EC vs Replication: {metric_titles[metric]} (VU:
{vu})")
    ax.set_xlabel('Bandwidth & Payload Size')
    ax.set_ylabel(metric_titles[metric])
    ax.tick_params(axis='x', rotation=90)

# single legend for all subplots
handles, labels = axes[0].get_legend_handles_labels()
fig.legend(handles, labels, title='System', loc='upper right')
fig.suptitle(f"EC vs Replication Metrics for {vu} Virtual Users",
fontsize=16)
fig.tight_layout(rect=[0, 0, 1, 0.96])
plt.show()

```



```
plot_by_vu()
```

```
/tmp/ipykernel_288815/1253657513.py:30: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_288815/1253657513.py:30: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_288815/1253657513.py:30: FutureWarning:
```


The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_288815/1253657513.py:30: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  

```



```

[31]: # 4. Create heatmaps to visualize the performance difference between EC and
      ↪Replication
      # This can help identify which configurations benefit most from EC vs
      ↪Replication

def plot_heatmaps_by_vu():
    df_ec = df[df['system'] == 'EC'].copy()
    df_repl = df[df['system'] == 'Replication'].copy()
    common_conditions = pd.merge(
        df_ec[['payload_size', 'virtual_user', 'bandwidth', 'bandwidth_num']],
        df_repl[['payload_size', 'virtual_user', 'bandwidth', 'bandwidth_num']],
        on=['payload_size', 'virtual_user', 'bandwidth', 'bandwidth_num']
    )
    if common_conditions.empty:
        return

    df_ec_f = pd.merge(df_ec, common_conditions,
        ↪on=['payload_size', 'virtual_user', 'bandwidth', 'bandwidth_num'])
    df_repl_f = pd.merge(df_repl, common_conditions,
        ↪on=['payload_size', 'virtual_user', 'bandwidth', 'bandwidth_num'])

    for vu in sorted(common_conditions['virtual_user'].unique()):
        fig, axes = plt.subplots(2, 2, figsize=(16, 12))
        axes = axes.flatten()
        for ax, metric in zip(axes, metrics):
            diff_data = []
            df_ec_vu = df_ec_f[df_ec_f['virtual_user'] == vu]
            df_repl_vu = df_repl_f[df_repl_f['virtual_user'] == vu]
            for _, ec_row in df_ec_vu.iterrows():
                repl_row = df_repl_vu[
                    (df_repl_vu['payload_size'] == ec_row['payload_size']) &
                    (df_repl_vu['bandwidth'] == ec_row['bandwidth'])
                ].iloc[0]

                if metric in ['med', 'p90', 'avg']:
                    denom = repl_row[metric] if repl_row[metric] != 0 else 1
                    diff = (ec_row[metric] - repl_row[metric]) / denom * 100
                else:
                    denom = ec_row[metric] if ec_row[metric] != 0 else 1
                    diff = (repl_row[metric] - ec_row[metric]) / denom * 100

            diff_data.append({
                'bandwidth': ec_row['bandwidth'],
                'bandwidth_num': ec_row['bandwidth_num'],
                'payload_size': ec_row['payload_size'],
                'diff': diff
            })

```

```

diff_df = pd.DataFrame(diff_data)
bw_order = (
    diff_df[['bandwidth', 'bandwidth_num']]
    .drop_duplicates()
    .sort_values('bandwidth_num')
    ['bandwidth']
    .tolist()
)
pivot = diff_df.pivot(index='bandwidth', columns='payload_size',
↪values='diff')\
    .reindex(bw_order)

sns.heatmap(pivot, annot=True, cmap='RdBu_r', center=0, fmt='.2f',
↪ax=ax)
ax.set_title(metric_titles[metric])
ax.set_xlabel('Payload Size (bytes)')
ax.set_ylabel('Bandwidth')

fig.suptitle(f'Performance Difference Heatmaps (VU: {vu})', fontsize=16)
plt.tight_layout(rect=[0,0,1,0.96])
plt.show()

plot_heatmaps_by_vu()

```

Performance Difference Heatmaps (VU: 1)

