

benchmark_k6_write

July 13, 2025

```
[8]: import json
import glob
import re
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Suffix for selecting benchmark results
suffix = "/write_avgnet/"

# Path to your benchmark results directory
results_dir = './results_store/_final/'
ec_pattern = results_dir + f'erasure{suffix}/_write_*.json'
replication_pattern = results_dir + f'replication{suffix}/_write_*.json'

# Helper to extract data from files
def extract_data(files, system_type):
    data = []
    for file in files:
        # Updated regex: matches kbit, mbit, gbit
        match = re.search(r'_write_(.+?)_(\d+)vu(?:_([0-9]+(?:kbit|mbit|gbit)))?'
↪'\.json', file)
        if not match:
            continue
        payload_size = match.group(1)
        if payload_size.endswith('b'):
            payload_size = payload_size[:-1]
        try:
            payload_size = int(payload_size)
        except ValueError:
            pass
        virtual_user = match.group(2)
        bandwidth = match.group(3) if match.group(3) else 'unlimited'
        with open(file) as f:
            j = json.load(f)
            # try summary.success_performance, else fallback to details.
↪http_req_duration
```

```

        sp = j.get('summary', {}).get('success_performance')
        if sp:
            med = sp.get('med', 0)
            p90 = sp.get('p(90)', 0)
            avg = sp.get('avg', 0)
        else:
            dur = j.get('details', {}).get('http_req_duration', {}).
↳get('values', {})
            med = dur.get('med', 0)
            p90 = dur.get('p(90)', 0)
            avg = dur.get('avg', 0)
        # extract request rate
        rate = j.get('summary', {}).get('reqs', {}).get('rate', 0)
        data.append({
            'system': system_type,
            'payload_size': payload_size,
            'virtual_user': int(virtual_user),
            'bandwidth': bandwidth,
            'med': med,
            'p90': p90,
            'avg': avg,
            'rate': rate
        })
    return data

def bandwidth_to_num(bw):
    if bw == 'unlimited':
        return float('inf')
    m = re.match(r'(\d+)(kbit|mbit|gbit)', bw)
    if not m:
        return float('inf')
    val, unit = int(m.group(1)), m.group(2)
    if unit == 'kbit':
        return val
    elif unit == 'mbit':
        return val * 1000
    elif unit == 'gbit':
        return val * 1000 * 1000
    return float('inf')

# Collect EC and Replication data
files_ec = glob.glob(ec_pattern)
files_replication = glob.glob(replication_pattern)
data = extract_data(files_ec, 'EC') + extract_data(files_replication,
↳'Replication')
df = pd.DataFrame(data)
if not df.empty:

```

```

# Add bandwidth_num for correct sorting
df['bandwidth_num'] = df['bandwidth'].apply(bandwidth_to_num)
# Sort by payload_size, virtual_user, bandwidth_num, then system for
↳grouped comparison
df = df.sort_values(['payload_size', 'virtual_user', 'bandwidth_num',
↳'system'])
df['combo'] = df.apply(lambda row:
↳f"{row['payload_size']}B_{row['virtual_user']}vu_{row['bandwidth']}", axis=1)
df.reset_index(drop=True, inplace=True)
df

```

```

[8]:
      system  payload_size  virtual_user  bandwidth      med  \
0         EC      200000           1      10mbit    699.992802
1  Replication      200000           1      10mbit   1073.896833
2         EC      200000           1      25mbit    341.796174
3  Replication      200000           1      25mbit    429.650143
4         EC      200000           1      40mbit    261.041856
5  Replication      200000           1      40mbit    268.192058
6         EC      200000           1      55mbit    236.579951
7  Replication      200000           1      55mbit    194.562242
8         EC      200000           1      70mbit    224.866249
9  Replication      200000           1      70mbit    152.595633
10        EC      400000           1      10mbit   1346.515693
11  Replication      400000           1      10mbit   2141.258066
12        EC      400000           1      25mbit    611.797533
13  Replication      400000           1      25mbit    856.185386
14        EC      400000           1      40mbit    441.182746
15  Replication      400000           1      40mbit    526.053664
16        EC      400000           1      55mbit    370.781581
17  Replication      400000           1      55mbit    380.759582
18        EC      400000           1      70mbit    335.465073
19  Replication      400000           1      70mbit    295.433864
20        EC      600000           1      10mbit   1994.064533
21  Replication      600000           1      10mbit   3207.928697
22        EC      600000           1      25mbit    874.594191
23  Replication      600000           1      25mbit   1279.382928
24        EC      600000           1      40mbit    615.983084
25  Replication      600000           1      40mbit    800.167324
26        EC      600000           1      55mbit    507.199511
27  Replication      600000           1      55mbit    569.969124
28        EC      600000           1      70mbit    458.925299
29  Replication      600000           1      70mbit    443.692196
30        EC      800000           1      10mbit   2640.590774
31  Replication      800000           1      10mbit   4575.859446
32        EC      800000           1      25mbit   1138.658848
33  Replication      800000           1      25mbit   1682.151866
34        EC      800000           1      40mbit    781.503670

```

35	Replication	800000	1	40mbit	1035.841426
36	EC	800000	1	55mbit	628.045337
37	Replication	800000	1	55mbit	754.886365
38	EC	800000	1	70mbit	544.997969
39	Replication	800000	1	70mbit	581.446226
40	EC	1000000	1	10mbit	3281.701913
41	Replication	1000000	1	10mbit	4126.683997
42	EC	1000000	1	25mbit	1394.745454
43	Replication	1000000	1	25mbit	2100.551557
44	EC	1000000	1	40mbit	969.631511
45	Replication	1000000	1	40mbit	1298.559282
46	EC	1000000	1	55mbit	760.338166
47	Replication	1000000	1	55mbit	940.436655
48	EC	1000000	1	70mbit	656.293743
49	Replication	1000000	1	70mbit	722.669057

	p90	avg	rate	bandwidth_num	combo
0	723.726307	701.948700	1.409262	10000	200000B_1vu_10mbit
1	1086.421647	1065.650607	0.931688	10000	200000B_1vu_10mbit
2	399.206527	347.620375	2.816840	25000	200000B_1vu_25mbit
3	466.753251	427.348248	2.297152	25000	200000B_1vu_25mbit
4	360.153438	275.510498	3.531695	40000	200000B_1vu_40mbit
5	289.197103	266.721019	3.639365	40000	200000B_1vu_40mbit
6	353.503377	253.428548	3.829072	55000	200000B_1vu_55mbit
7	207.905265	193.301509	4.969445	55000	200000B_1vu_55mbit
8	365.015640	241.722793	4.010621	70000	200000B_1vu_70mbit
9	163.087706	151.345583	6.277621	70000	200000B_1vu_70mbit
10	1380.522840	1349.873754	0.733990	10000	400000B_1vu_10mbit
11	2299.629240	2110.231205	0.470518	10000	400000B_1vu_10mbit
12	672.821147	616.714089	1.590676	25000	400000B_1vu_25mbit
13	982.572614	846.049949	1.166899	25000	400000B_1vu_25mbit
14	547.281528	457.927517	2.138820	40000	400000B_1vu_40mbit
15	594.218360	526.925748	1.862124	40000	400000B_1vu_40mbit
16	566.006723	402.389760	2.422453	55000	400000B_1vu_55mbit
17	415.795460	381.054784	2.559696	55000	400000B_1vu_55mbit
18	572.220025	374.013042	2.604556	70000	400000B_1vu_70mbit
19	325.836070	296.484580	3.253635	70000	400000B_1vu_70mbit
20	2021.718340	1992.436063	0.496600	10000	600000B_1vu_10mbit
21	3333.020765	3140.549373	0.316151	10000	600000B_1vu_10mbit
22	933.686783	879.782181	1.116080	25000	600000B_1vu_25mbit
23	1431.338001	1265.027102	0.778043	25000	600000B_1vu_25mbit
24	771.424799	636.473482	1.532572	40000	600000B_1vu_40mbit
25	891.451074	790.580834	1.241447	40000	600000B_1vu_40mbit
26	735.781840	543.572890	1.790709	55000	600000B_1vu_55mbit
27	665.196284	570.356202	1.706402	55000	600000B_1vu_55mbit
28	720.457737	496.046357	1.953569	70000	600000B_1vu_70mbit
29	518.723513	445.896458	2.171622	70000	600000B_1vu_70mbit

30	2676.778082	2640.417792	0.374630	10000	800000B_1vu_10mbit
31	10689.518645	5703.423259	0.174301	10000	800000B_1vu_10mbit
32	1214.315463	1145.674224	0.852595	25000	800000B_1vu_25mbit
33	1983.821932	1671.141309	0.588509	25000	800000B_1vu_25mbit
34	930.897583	805.201959	1.195295	40000	800000B_1vu_40mbit
35	1282.618340	1035.320619	0.938890	40000	800000B_1vu_40mbit
36	845.940564	669.386850	1.433106	55000	800000B_1vu_55mbit
37	873.146131	747.611853	1.288427	55000	800000B_1vu_55mbit
38	811.969817	598.569975	1.594321	70000	800000B_1vu_70mbit
39	699.595252	579.789730	1.639239	70000	800000B_1vu_70mbit
40	3309.682499	3282.157200	0.301631	10000	1000000B_1vu_10mbit
41	7102.105745	6856.459491	0.145046	10000	1000000B_1vu_10mbit
42	1467.124810	1401.247301	0.696896	25000	1000000B_1vu_25mbit
43	2584.093762	2081.660721	0.471932	25000	1000000B_1vu_25mbit
44	1092.546838	975.910794	0.987347	40000	1000000B_1vu_40mbit
45	1617.774552	1287.491610	0.755246	40000	1000000B_1vu_40mbit
46	947.808956	787.650352	1.212604	55000	1000000B_1vu_55mbit
47	1161.139322	932.183974	1.032365	55000	1000000B_1vu_55mbit
48	903.999017	701.980075	1.354233	70000	1000000B_1vu_70mbit
49	914.024479	722.928048	1.316559	70000	1000000B_1vu_70mbit

```
[9]: # Plotting: Compare EC vs Replication for each metric, grouped by
      ↪ (payload_size, virtual_user)
```

```
sns.set_style("whitegrid")
metrics = ['med', 'p90', 'avg', 'rate']
metric_titles = {
    'med': 'Median Latency (ms)',
    'p90': 'P90 Latency (ms)',
    'avg': 'Average Latency (ms)',
    'rate': 'Request Rate (req/s)'
}

fig, axes = plt.subplots(2, 2, figsize=(20, 12))
axes = axes.flatten()

for ax, metric in zip(axes, metrics):
    sns.barplot(
        data=df,
        x='combo',
        y=metric,
        hue='system',
        ci=None,
        dodge=True,
        ax=ax
    )
    ax.set_title(f'EC vs Replication: {metric_titles[metric]}')
    ax.set_xlabel('Payload Size & Virtual Users')
```

```

ax.set_ylabel(metric_titles[metric])
ax.tick_params(axis='x', rotation=90)

# single legend for all subplots
handles, labels = axes[0].get_legend_handles_labels()
fig.legend(handles, labels, title='System', loc='upper right')

fig.tight_layout()
plt.show()

```

/tmp/ipykernel_171471/1184367366.py:15: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```

sns.barplot(
/tmp/ipykernel_171471/1184367366.py:15: FutureWarning:

```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```

sns.barplot(
/tmp/ipykernel_171471/1184367366.py:15: FutureWarning:

```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```

sns.barplot(
/tmp/ipykernel_171471/1184367366.py:15: FutureWarning:


```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```

sns.barplot(

```



```
[10]: # Create grouped visualizations for bandwidth, payload size, and virtual users
```

```
# First, let's check what distinct values we have for each parameter
print(f"Unique bandwidth values: {df['bandwidth'].unique()}")
print(f"Unique payload sizes: {df['payload_size'].unique()}")
print(f"Unique virtual users: {df['virtual_user'].unique()}")
```

```
Unique bandwidth values: ['10mbit' '25mbit' '40mbit' '55mbit' '70mbit']
Unique payload sizes: [200000 400000 600000 800000 1000000]
Unique virtual users: [1]
```

```
[11]: # 1. Grouping by bandwidth
```

```
def plot_by_bandwidth():
    # Sort bandwidths numerically using bandwidth_num
    bw_order = (
        df[['bandwidth', 'bandwidth_num']]
        .drop_duplicates()
        .sort_values('bandwidth_num')
        .bandwidth.tolist()
    )
    for bw in bw_order:
        df_bw = df[df['bandwidth'] == bw]
        if df_bw.empty:
            continue

    # set up 2 x 2 subplots for the four metrics
```

```

fig, axes = plt.subplots(2, 2, figsize=(20, 12))
axes = axes.flatten()
for ax, metric in zip(axes, metrics):
    # label for x-axis
    df_bw['payload_vu'] = df_bw.apply(
        lambda row: f"{row['payload_size']}B_{row['virtual_user']}vu",
↪axis=1
    )
    sns.barplot(
        data=df_bw,
        x='payload_vu',
        y=metric,
        hue='system',
        ci=None,
        dodge=True,
        ax=ax
    )
    ax.set_title(f'EC vs Replication: {metric_titles[metric]} at {bw}
↪Bandwidth')
    ax.set_xlabel('Payload Size & Virtual Users')
    ax.set_ylabel(metric_titles[metric])
    ax.tick_params(axis='x', rotation=90)

    # single legend for all subplots
    handles, labels = axes[0].get_legend_handles_labels()
    fig.legend(handles, labels, title='System', loc='upper right')
    fig.suptitle(f'EC vs Replication Metrics at {bw} Bandwidth',
↪fontsize=16)
    fig.tight_layout(rect=[0, 0, 1, 0.97])
    plt.show()

plot_by_bandwidth()

```

/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
```


Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.


```
sns.barplot(
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(
```



```
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
```

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.


```
sns.barplot(
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(
```



```
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

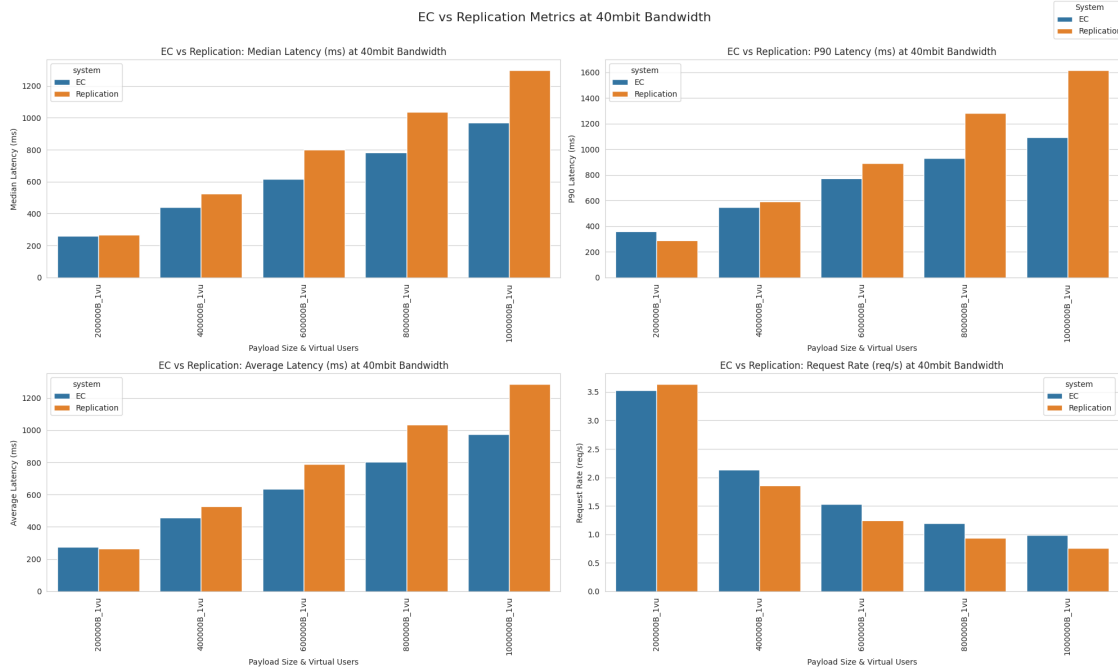
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  

```



```
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
```

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.


```
sns.barplot(
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(
```



```
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_171471/2927317137.py:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_bw['payload_vu'] = df_bw.apply(  
/tmp/ipykernel_171471/2927317137.py:23: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  

```



```
[12]: # 2. Grouping by payload size (all metrics in one 2 x 2 display)
def plot_by_payload():
    payloads = df['payload_size'].unique()
    for payload in payloads:
        df_payload = df[df['payload_size'] == payload]
        if df_payload.empty:
            continue

        # Sort bandwidths numerically for this payload
        bw_order = (
            df_payload[['bandwidth', 'bandwidth_num']]
            .drop_duplicates()
            .sort_values('bandwidth_num')
            .bandwidth.tolist()
        )
        # Create combined x-axis labels and enforce ordering
        df_payload['bw_vu'] = df_payload.apply(
            lambda row: f"{row['bandwidth']}_{row['virtual_user']}_vu", axis=1
        )
        vu_order = sorted(df_payload['virtual_user'].unique())
        x_order = [f"{bw}_{vu}_vu" for bw in bw_order for vu in vu_order]
        df_payload['bw_vu'] = pd.Categorical(df_payload['bw_vu'],
        categories=x_order, ordered=True)

        # set up 2 x 2 subplots for the four metrics
```



```

fig, axes = plt.subplots(2, 2, figsize=(20, 12))
axes = axes.flatten()
for ax, metric in zip(axes, metrics):
    sns.barplot(
        data=df_payload,
        x='bw_vu',
        y=metric,
        hue='system',
        ci=None,
        dodge=True,
        ax=ax
    )
    ax.set_title(f"EC vs Replication: {metric_titles[metric]} for {payload}B")
    ax.set_xlabel("Bandwidth & Virtual Users")
    ax.set_ylabel(metric_titles[metric])
    ax.tick_params(axis='x', rotation=90)

    # single legend for all subplots
    handles, labels = axes[0].get_legend_handles_labels()
    fig.legend(handles, labels, title='System', loc='upper right')
    fig.suptitle(f"EC vs Replication Metrics for {payload}B Payload Size",
    ↪fontsize=16)
    fig.tight_layout(rect=[0, 0, 1, 0.96])
    plt.show()

plot_by_payload()

```

/tmp/ipykernel_171471/1240405451.py:17: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = df_payload.apply(
```

/tmp/ipykernel_171471/1240405451.py:22: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = pd.Categorical(df_payload['bw_vu'], categories=x_order,
ordered=True)
```

/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

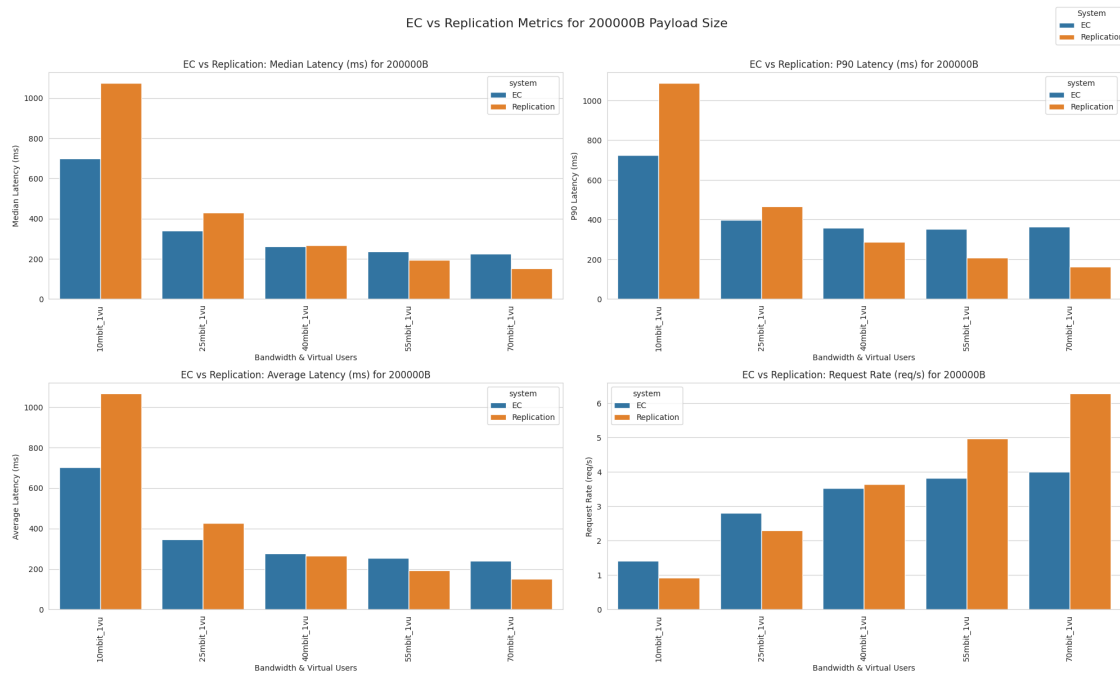
```
sns.barplot(
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
```



```
/tmp/ipykernel_171471/1240405451.py:17: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = df_payload.apply(
/tmp/ipykernel_171471/1240405451.py:22: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
```

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = pd.Categorical(df_payload['bw_vu'], categories=x_order,
ordered=True)
```

/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.


```
sns.barplot(
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The ``ci`` parameter is deprecated. Use ``errorbar=None`` for the same effect.

```
sns.barplot(
```



```
/tmp/ipykernel_171471/1240405451.py:17: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = df_payload.apply(  
/tmp/ipykernel_171471/1240405451.py:22: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = pd.Categorical(df_payload['bw_vu'], categories=x_order,  
ordered=True)  
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```


The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  

```



```
/tmp/ipykernel_171471/1240405451.py:17: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = df_payload.apply(
/tmp/ipykernel_171471/1240405451.py:22: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = pd.Categorical(df_payload['bw_vu'], categories=x_order,
ordered=True)
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```


The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  

```



```
/tmp/ipykernel_171471/1240405451.py:17: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = df_payload.apply(  
/tmp/ipykernel_171471/1240405451.py:22: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_payload['bw_vu'] = pd.Categorical(df_payload['bw_vu'], categories=x_order,  
ordered=True)  
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

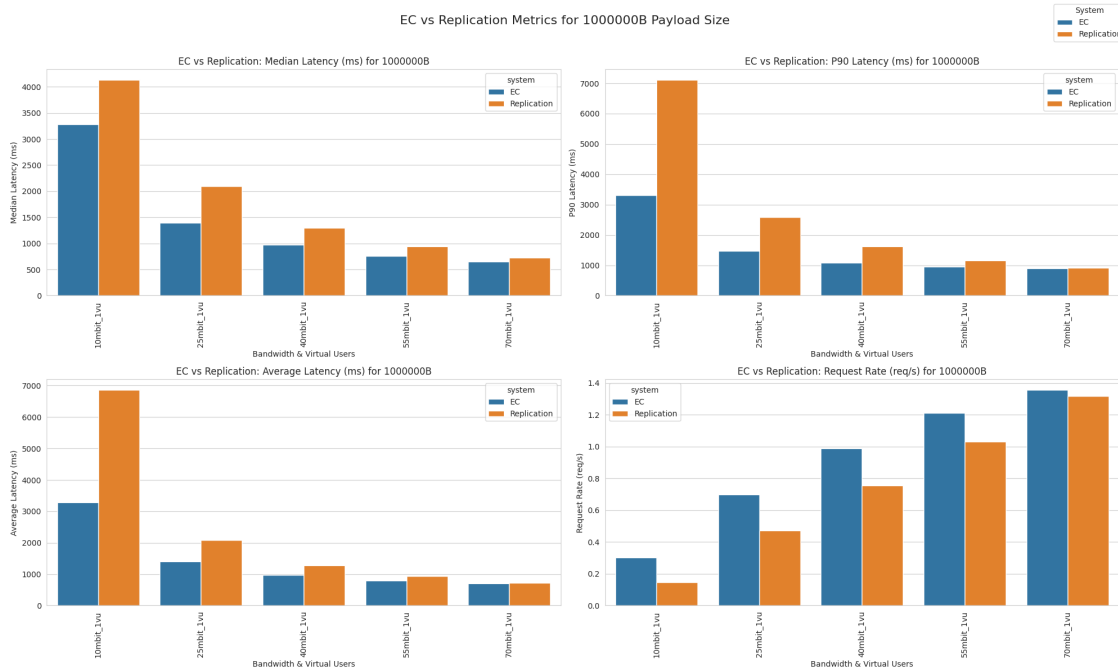
```
sns.barplot(
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
/tmp/ipykernel_171471/1240405451.py:28: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(
```



```
[13]: # 3. Grouping by virtual users (all metrics in one 2 x 2 display)
def plot_by_vu():
    vus = df['virtual_user'].unique()
    for vu in vus:
        df_vu = df[df['virtual_user'] == vu]
```

```

if df_vu.empty:
    continue

# set up 2 x 2 subplots for the four metrics
fig, axes = plt.subplots(2, 2, figsize=(20, 12))
axes = axes.flatten()

# precompute ordering for x-axis
bw_order = (
    df_vu[['bandwidth', 'bandwidth_num']]
    .drop_duplicates()
    .sort_values('bandwidth_num')
    .bandwidth.tolist()
)
ps_order = sorted(df_vu['payload_size'].unique())
x_order = [f"{bw}_{ps}B" for bw in bw_order for ps in ps_order]

# apply categorical ordering
df_vu['bw_payload'] = df_vu.apply(
    lambda row: f"{row['bandwidth']}_{row['payload_size']}B", axis=1
)
df_vu['bw_payload'] = pd.Categorical(df_vu['bw_payload'],
categories=x_order, ordered=True)

for ax, metric in zip(axes, metrics):
    sns.barplot(
        data=df_vu,
        x='bw_payload',
        y=metric,
        hue='system',
        ci=None,
        dodge=True,
        ax=ax
    )
    ax.set_title(f"EC vs Replication: {metric_titles[metric]} (VU: {vu})")
    ax.set_xlabel('Bandwidth & Payload Size')
    ax.set_ylabel(metric_titles[metric])
    ax.tick_params(axis='x', rotation=90)

# single legend for all subplots
handles, labels = axes[0].get_legend_handles_labels()
fig.legend(handles, labels, title='System', loc='upper right')
fig.suptitle(f"EC vs Replication Metrics for {vu} Virtual Users",
fontsize=16)
fig.tight_layout(rect=[0, 0, 1, 0.96])
plt.show()

```



```
plot_by_vu()
```

```
/tmp/ipykernel_171471/1253657513.py:30: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_171471/1253657513.py:30: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_171471/1253657513.py:30: FutureWarning:
```


The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  
/tmp/ipykernel_171471/1253657513.py:30: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(  

```



```

[14]: # 4. Create heatmaps to visualize the performance difference between EC and
      ↪Replication
      # This can help identify which configurations benefit most from EC vs
      ↪Replication

def plot_heatmaps_by_vu():
    df_ec = df[df['system'] == 'EC'].copy()
    df_repl = df[df['system'] == 'Replication'].copy()
    common_conditions = pd.merge(
        df_ec[['payload_size', 'virtual_user', 'bandwidth', 'bandwidth_num']],
        df_repl[['payload_size', 'virtual_user', 'bandwidth', 'bandwidth_num']],
        on=['payload_size', 'virtual_user', 'bandwidth', 'bandwidth_num']
    )
    if common_conditions.empty:
        return

    df_ec_f = pd.merge(df_ec, common_conditions,
        ↪on=['payload_size', 'virtual_user', 'bandwidth', 'bandwidth_num'])
    df_repl_f = pd.merge(df_repl, common_conditions,
        ↪on=['payload_size', 'virtual_user', 'bandwidth', 'bandwidth_num'])

    for vu in sorted(common_conditions['virtual_user'].unique()):
        fig, axes = plt.subplots(2, 2, figsize=(16, 12))
        axes = axes.flatten()
        for ax, metric in zip(axes, metrics):
            diff_data = []
            df_ec_vu = df_ec_f[df_ec_f['virtual_user'] == vu]
            df_repl_vu = df_repl_f[df_repl_f['virtual_user'] == vu]
            for _, ec_row in df_ec_vu.iterrows():
                repl_row = df_repl_vu[
                    (df_repl_vu['payload_size'] == ec_row['payload_size']) &
                    (df_repl_vu['bandwidth'] == ec_row['bandwidth'])
                ].iloc[0]

                if metric in ['med', 'p90', 'avg']:
                    denom = repl_row[metric] if repl_row[metric] != 0 else 1
                    diff = (ec_row[metric] - repl_row[metric]) / denom * 100
                else:
                    denom = ec_row[metric] if ec_row[metric] != 0 else 1
                    diff = (repl_row[metric] - ec_row[metric]) / denom * 100

                diff_data.append({
                    'bandwidth': ec_row['bandwidth'],
                    'bandwidth_num': ec_row['bandwidth_num'],
                    'payload_size': ec_row['payload_size'],
                    'diff': diff
                })

```

```

diff_df = pd.DataFrame(diff_data)
bw_order = (
    diff_df[['bandwidth', 'bandwidth_num']]
    .drop_duplicates()
    .sort_values('bandwidth_num')
    ['bandwidth']
    .tolist()
)
pivot = diff_df.pivot(index='bandwidth', columns='payload_size',
↪values='diff')\
    .reindex(bw_order)

sns.heatmap(pivot, annot=True, cmap='RdBu_r', center=0, fmt='.2f',
↪ax=ax)
ax.set_title(metric_titles[metric])
ax.set_xlabel('Payload Size (bytes)')
ax.set_ylabel('Bandwidth')

fig.suptitle(f'Performance Difference Heatmaps (VU: {vu})', fontsize=16)
plt.tight_layout(rect=[0,0,1,0.96])
plt.show()

plot_heatmaps_by_vu()

```

Performance Difference Heatmaps (VU: 1)

